


Graphical Abstract

Adsorption ML needs source-aware validation



High interpolation accuracy, applicability domain, and external generalization are different scientific claims.

Highlights

- Source-specific adsorption models reached R^2 up to 0.992.
- Sparse neutral-low-concentration cases were excluded from primary modeling.
- Source-origin heterogeneity explained validation sensitivity.
- Row-wise validation overstated source-transfer performance.
- Reference-held-out sensitivity reached $R^2 = 0.330$ in the main domain.

Source-aware machine learning for AMD-relevant metal adsorption: distinguishing high-accuracy interpolation from cross-source generalization

Yash Bisht^a, Margreth Tadie^b, Susmit Chitransh^b

^a*Department of Mechanical and Industrial Engineering, Indian Institute of Technology Roorkee, Roorkee 247667, Uttarakhand, India*

^b*Department of Chemical Engineering, Stellenbosch University, Stellenbosch, Cape Town, South Africa*

Abstract

Acid mine drainage (AMD) is a persistent water-quality challenge because acidic, sulfate-rich drainage mobilizes metal and metalloid contaminants. Adsorption is widely investigated for AMD-relevant contaminant remediation, and machine learning (ML) can support rapid prediction of removal efficiency; however, reported accuracy can depend strongly on feature choice, applicability-domain coverage, and validation design. This study develops a source-aware and leakage-aware validation framework for adsorption-based metal-contaminant removal-efficiency prediction using a literature-compiled experimental dataset. The contribution is an applicability-domain audit and validation framework, not a claim of a universally deployable adsorption predictor. Source awareness was introduced by tracking the literature reference from which each adsorption record originated. Leakage awareness was addressed by testing whether model performance remained reliable when familiar rows, adsorbents, or entire literature sources were removed from training, and by separating capacity-inclusive process-monitoring models from capacity-free pre-experiment screening mod-

els. The raw dataset contained 566 adsorption records; after deterministic numeric parsing and quality filtering, 563 valid records remained. Because the neutral/low-concentration regime contained only 26 records, primary modeling was restricted to a 537-record domain covering acidic/high-concentration, acidic/low-concentration, and neutral/high-concentration conditions. Track A quantified source-specific interpolation through within-source cross-validation, where XGBoost reached R^2 values of 0.992, 0.974, and 0.960 for three major source groups. Track B tested broader generalization using random row-wise, adsorbent-held-out, and reference-held-out splits. In the restricted primary domain, the full capacity-inclusive LightGBM model achieved $R^2 = 0.806$ under random row-wise validation. Under stricter validation, the best adsorbent-held-out performance was obtained by a capacity-free numeric LightGBM model ($R^2 = 0.532$), while reference-held-out performance reached $R^2 = 0.330$ for a capacity-inclusive LightGBM model and $R^2 = 0.305$ for the best capacity-free model. A source-domain audit showed that source-origin heterogeneity strongly affected validation sensitivity. These results show that adsorption-ML accuracy can be heavily validation-dependent: high accuracy may indicate interpolation within familiar sources, whereas lower reference-held-out accuracy diagnoses the present gap in external-source transfer.

Keywords: acid mine drainage, adsorption, metal removal efficiency, machine learning, XGBoost, LightGBM, source-aware validation

1. Introduction

Mining activities are a major contributor to water-quality degradation, particularly through the generation of acid mine drainage (AMD). AMD is produced

mainly through oxidative dissolution of sulfide minerals and is commonly associated with low pH, high sulfate concentration, and elevated dissolved metal or metalloid concentrations [1]. These conditions can mobilize contaminants such as Fe, Mn, Pb, Cu, Zn, Cr, As, Cd, and Ni, creating long-term risks for receiving water bodies, soils, aquatic ecosystems, and downstream water reuse. Because AMD chemistry varies across mine sites, ore bodies, weathering history, and treatment conditions, robust prediction of contaminant removal remains a difficult environmental-chemical-engineering problem.

Adsorption is widely investigated for AMD-relevant contaminant remediation because it can operate under mild conditions, can use low-cost or waste-derived adsorbents, and can be optimized through pH, adsorbent dose, contact time, temperature, surface chemistry, and initial contaminant concentration. Recent reviews emphasize adsorption as an active research direction for AMD treatment, including carbonaceous, mineral, oxide, biomass-derived, and composite adsorbents [2]. The expansion of experimental adsorption datasets has therefore encouraged the use of machine learning (ML), a subset of artificial intelligence, to predict removal efficiency and accelerate adsorbent or operating-condition screening.

ML is attractive in this setting because adsorption performance depends on nonlinear interactions among solution chemistry, contact time, adsorbent dose, contaminant identity, material surface area, and source-specific synthesis or reporting practices. Data-driven models can capture multivariate trends that are difficult to express with a single empirical equation, and they can provide rapid screening support once trained on reliable data. However, the validity of an ML model depends not only on the algorithm but also on whether the test data are genuinely independent of the training data and whether the input features are available

at the intended stage of deployment.

30 Recent adsorption and environmental-remediation ML studies demonstrate that high predictive accuracy is possible (Table 1). Hafsa et al. modeled heavy-metal adsorption efficiencies using multiple ML algorithms and reported R^2 values in the 0.96–0.99 range under ten-fold cross-validation, with experimental data supplemented by synthetic interpolation [3]. Lu et al. predicted
35 heavy-metal removal by chitosan-based flocculants using random forests and reported $R^2 = 0.9354$, identifying solution pH and flocculant molecular weight as important drivers [4]. Wang et al. used interpretable ML for heavy-metal removal by biochar and reported XGBoost test performance near $R^2 = 0.99$, together with SHAP and partial-dependence analyses to relate removal efficiency to biochar and
40 operating variables [5, 8]. Recent ensemble-learning work on biochar adsorption similarly reported strong XGBoost performance for heavy-metal adsorption efficiency prediction [6]. Within the scope of *Journal of Environmental Chemical Engineering*, interpretable XGBoost modeling has also been applied to heavy-metal removal efficiency in electrokinetic remediation, demonstrating the journal
45 relevance of transparent ML workflows for environmental treatment systems [7].

 The central challenge is that high validation accuracy does not necessarily imply reliable generalization to unseen laboratories, adsorbents, contaminant systems, or literature sources. In many adsorption ML studies, performance is estimated using random row-wise validation, where individual experimental records
50 are randomly divided between training and testing. This can be useful for measuring interpolation within a compiled dataset, but it may also place closely related records from the same study dataset, adsorbent series, or operating campaign in both the training and test sets. As a result, the test set can appear independent while

Table 1: Representative adsorption/remediation ML studies and validation gaps motivating the present work.

Study/domain	Reported strength	Gap addressed here
Heavy-metal adsorption ML [3]	High cross-validation R^2 values, including synthetic interpolation support	Cross-validation can indicate interpolation performance but may not test transfer to unseen literature sources or adsorbent campaigns.
Chitosan-based flocculant removal [4]	Random-forest prediction with $R^2 > 0.93$ and interpretable pH/molecular-weight effects	Strong within-domain prediction does not by itself establish source-held-out or material-held-out generalization.
Biochar heavy-metal adsorption [5, 6]	XGBoost/ensemble models with high test performance and interpretable feature effects	Adsorbent families, feedstocks, and protocols may appear in both train and test sets under row-wise validation.
Electrokinetic remediation ML in JECE [7]	Demonstrates journal relevance of explainable ML for contaminant-removal efficiency	Broader adsorption workflows still need explicit feature-use-case separation and source-aware validation.
PFAS adsorption/partitioning ML [11]	Random-split performance substantially exceeded leave-compound-out transfer	Supports the need to distinguish interpolation from transfer to unseen chemical/source domains.

still containing familiar source-specific patterns. In ML-based science, such data
 55 leakage and validation mismatch can produce overoptimistic estimates of deploy-
 able performance [9]. Environmental ML best-practice guidance similarly em-
 phasizes careful data splitting, leakage management, randomness assessment, and
 evaluation design [10]. These concerns are particularly important for adsorption
 datasets because different studies may use different adsorbent synthesis routes,
 60 contaminant species, concentration ranges, analytical protocols, and reporting con-
 ventions.

Feature definition can further blur the interpretation of reported accuracy.
 Adsorption capacity and removal efficiency are both derived from concentration
 changes in many batch adsorption experiments:

$$\text{RE}(\%) = 100 \frac{C_0 - C_e}{C_0}, \quad (1)$$

$$q_e = \frac{(C_0 - C_e)V}{m}, \quad (2)$$

65 where C_0 and C_e are the initial and equilibrium concentrations, V is the solu-
 tion volume, and m is the adsorbent mass. Because both RE and q_e depend on
 the concentration change $C_0 - C_e$, adsorption capacity can carry strong informa-
 tion about removal efficiency. A model that uses capacity is therefore useful for
 process monitoring or post-experiment analysis, where capacity has already been
 70 measured or estimated. However, the same model should be treated cautiously for
 pre-experiment screening, because the capacity of a candidate adsorbent under a
 new condition is usually not known before the experiment is performed. Therefore,
 capacity-inclusive and capacity-free models answer different practical questions
 and should not be interpreted in the same way. Recent PFAS adsorption model-

75 ing has similarly shown that random-split performance can be much stronger than
leave-one-compound-out performance, illustrating that interpolation and transfer
to unseen chemical domains are distinct tasks [11].

This study addresses these issues by developing a source-aware and leakage-
aware validation framework for AMD-relevant metal-contaminant adsorption
80 removal-efficiency prediction across a broad literature-derived dataset spanning
189 normalized source references, 219 adsorbents, and 16 adsorbate classes.
Here, source awareness means tracking the literature reference from which each
adsorption record originated, rather than treating all rows as fully independent
observations. Leakage awareness means testing whether model performance re-
85 mains reliable when familiar information is removed from training, and separating
feature sets that may be appropriate after an experiment from feature sets that are
suitable before an experiment. The framework is intended to support ML-based
mathematical modeling for understanding the combined and individual effects of
adsorption variables while avoiding unsupported claims of universal deployment.

90 The framework separates two aspects that are often mixed in adsorption ML
studies: what information is available to the model, and how independent the test
data are from the training data. Four feature sets were constructed to represent
different levels of information availability. A full capacity-inclusive feature set
included adsorption capacity and capacity-derived predictors and was interpreted
95 as a process-monitoring or post-experiment analysis model. Three capacity-free
feature sets excluded adsorption capacity and were interpreted as pre-experiment
screening models, because the capacity of a candidate adsorbent is usually not
known before testing. Among the capacity-free models, one represented adsor-
bates using ion descriptors, one used normalized adsorbate identity as a categori-

100 cal label, and one used only numeric operating, material, and engineered process variables as a strict screening baseline.

Two validation tracks were then used to separate interpolation from broader generalization. Track A evaluated source-specific interpolation by training and testing within individual literature sources. This asks whether high R^2 values com-
105 parable to those reported in the adsorption ML literature can be reproduced inside known source groups. Track B evaluated broader generalization using increasingly strict split designs. In random row-wise validation, individual records were randomly divided between training and testing. In adsorbent-held-out validation, selected adsorbents were kept completely unseen during training and used only
110 for testing. In reference-held-out validation, all records from selected literature sources were kept unseen during training and used only for testing. Thus, the study asks three practical questions: random validation asks whether the model can interpolate inside the compiled dataset; adsorbent-held-out validation asks whether it can transfer to unseen material names; and reference-held-out validation asks
115 whether it can transfer to unseen literature or source domains. The novelty of this study is therefore not simply a new regression model, but a source-aware and leakage-aware reporting framework that clarifies when high R^2 reflects valid interpolation within known experimental domains and when additional validation is needed before claiming deployment to new adsorbents, studies, or laboratories.

120 **2. Materials and methods**

2.1. Dataset and target variable

The dataset was compiled from literature-reported adsorption experiments for AMD-relevant metal and metalloid contaminants. Each record corresponds to one

adsorption condition with a reported removal efficiency. The raw dataset contained 566 rows and 13 columns. After numeric parsing, missing-value filtering, and constraining removal efficiency to the physically meaningful range 0–100%, 563 records remained. The cleaned dataset included 189 normalized source references, 219 adsorbents, and 16 normalized adsorbate classes. The primary modeling domain excluded the sparse neutral/low-concentration regime, leaving 537 records for model training and testing.

The dataset was intended as a heterogeneous multi-source benchmark for evaluating validation design in adsorption ML, rather than as a systematic meta-analysis of all AMD treatment or heavy-metal adsorption studies. Records were retained when removal efficiency and the required operating or material variables were reported or could be deterministically parsed from tabulated values. Records were excluded when essential numeric variables were missing, when removal efficiency could not be constrained to the physical range of 0–100%, or when adsorbent, adsorbate, or source identity could not be normalized consistently.

The target variable was removal efficiency, RE (%). Input variables included adsorbent name, adsorbate identity, surface area, adsorption capacity, initial concentration, contact time, adsorbent dose, agitation speed, initial pH, temperature, and source reference. The source identifier, denoted Ref_group, was used as a source-origin proxy because explicit laboratory-origin labels were not available for all records. The pH-concentration regime label was not used as a model target or as a model input feature; it was created only for stratified splitting, diagnostic reporting, and applicability-domain definition.

To interpret validation sensitivity, a source-domain audit was performed on the restricted primary domain. This audit quantified the number of records per nor-

malized source, the number of singleton sources, the fraction of rows contributed
 150 by the largest source groups, source-wise RE means and standard deviations, ad-
 sorbent and adsorbate diversity within each source, and pH-concentration-regime
 coverage. A one-way source-label variance diagnostic was also computed as

$$\eta_{\text{source}}^2 = \frac{\sum_g n_g (\bar{y}_g - \bar{y})^2}{\sum_i (y_i - \bar{y})^2}, \quad (3)$$

where g indexes source groups, n_g is the number of records in source g , $\bar{y}_g$ is the
 mean RE for source g , and $\bar{y}$ is the overall mean RE. This statistic was used only
 155 to diagnose source imprint in the dataset; Ref_group was not used as a predictor
 in the deployed Track B models.

2.2. Numeric cleaning and categorical normalization

Literature-derived adsorption tables often contain nonuniform numeric format-
 ting. Numeric entries were cleaned using deterministic rules: decimal commas
 160 were converted to decimal points, uncertainty and approximate symbols were re-
 moved, multiple values in one cell were reduced to the first numeric value, and
 nonnumeric text was parsed using regular expressions. Rows with missing values
 in required numeric variables were removed. Adsorbate names were normalized
 to reduce duplicates such as elemental and valence-state aliases. Adsorbent names
 165 were grouped into broad families, including carbon-based, mineral/oxide, poly-
 mer, biomass, and composite/other categories.

The cleaned data were partitioned into four pH-concentration regimes using
 pH = 7 as the acidic/near-neutral boundary and the dataset median initial concen-
 tration, $C_0 = 50$ ppm, as the low/high concentration boundary (Table 2). The
 170 median C_0 is the middle value of the initial-concentration distribution and pro-
 vides a data-driven threshold for separating lower- and higher-concentration ex-

periments. These regimes were not used as prediction targets and were not passed to XGBoost, LightGBM, or ExtraTrees as predictor variables. Instead, the models predicted RE (%) from the original experimental and material variables, including numerical pH and numerical initial concentration. Regime labels were used only for stratified random splitting and diagnostic analysis, ensuring that train-test splits preserved chemically meaningful coverage across major operating regions. The acidic regimes dominated the dataset, whereas the neutral/low- C_0 regime contained only 26 observations. Because this regime was too small for stable split-wise evaluation, it was excluded from the primary modeling domain and treated as outside the present applicability domain. The excluded regime was not discarded from the dataset audit; it was reported explicitly as an out-of-domain limitation to avoid presenting unstable validation results from a poorly represented operating region.

Table 2: pH-concentration regime distribution in the cleaned dataset.

Regime	Count	Mean pH	Mean C_0 (ppm)	Mean RE (%)
Acidic_HighC0	246	4.51	190.27	70.30
Acidic_LowC0	231	5.12	11.04	74.06
Neutral_HighC0	60	7.65	166.51	68.05
Neutral_LowC0	26	7.47	20.24	81.41

Neutral_LowC0 was excluded from primary modeling because it contained only 26 observations; the other three regimes formed the 537-record primary domain.

2.3. Feature sets

Four feature sets were constructed to distinguish scientific use cases (Table 3). This separation is important because a model that includes adsorption capacity answers a different question from a model intended for pre-experiment adsor-




Figure 1: Composition of the cleaned adsorption dataset after deterministic numeric parsing and quality filtering. The sparse neutral/low- C_0 regime was retained in the dataset audit but excluded from the primary 537-record modeling domain.

bent screening. Adsorption capacity is closely related to removal efficiency because both are calculated from concentration change during adsorption. In most
190 batch adsorption studies, capacity is not known before an experiment; it is calculated after measuring or estimating equilibrium concentration. Capacity-inclusive models are therefore best interpreted as descriptive, process-monitoring, or post-experiment analysis models. In contrast, capacity-free models use information that is normally available before testing, such as pH, initial concentration, dose, contact time, temperature, surface area, adsorbent family, and adsorbate information.
195 The base experimental inputs were surface area, initial concentration, contact time, dose, agitation speed, initial pH, and temperature. The process-monitoring feature set additionally used adsorption capacity and capacity-derived terms, whereas the screening feature sets deliberately removed these terms.

200 Physics-inspired variables included site density, logarithmic contact time, mixing index, driving-force proxy, and acidity strength:

$$\text{Site Density} = \frac{\text{dose} \times \text{surface area}}{C_0}, \quad (4)$$

$$\text{LogTime} = \log(1 + \text{contact time}), \quad (5)$$

$$\text{Mixing Index} = \text{RPM} \times \text{LogTime}, \quad (6)$$

$$\text{Driving Force} = \frac{C_0}{\text{dose}}, \quad (7)$$

$$\text{Acidity Strength} = |\text{pH} - 7|. \quad (8)$$

These engineered variables were used as heuristic process descriptors rather than mechanistic adsorption-law parameters. They expose the tree models to simple combinations of dose, surface area, concentration, contact time, agitation, and
205 acidity that are commonly relevant to adsorption behavior.

Table 3: Feature-set definitions used to separate process-monitoring and screening tasks.

Feature set	Included information	Intended use
Full capacity process model	Operating variables, surface area, adsorption capacity, capacity-derived predictors, engineered process variables, ion descriptors, adsorbent family, and adsorbate identity	Descriptive modeling, process monitoring, or post-experiment analysis
Screening-safe plus ion	Operating variables, surface area, engineered process variables, ion descriptors, and adsorbent family; excludes adsorption capacity and capacity-derived terms	Pre-experiment screening with basic ion chemistry
Screening-safe plus adsorbate one-hot	Operating variables, surface area, engineered process variables, normalized adsorbate identity, and adsorbent family; excludes adsorption capacity and capacity-derived terms	Pre-experiment screening when contaminant identity is known
Screening-safe numeric-only	Operating variables, surface area, and engineered process variables; excludes adsorption capacity, capacity-derived terms, and direct adsorbate identity	Strict screening baseline

Ion descriptors were added to represent adsorbates through simple aqueous-chemistry properties rather than only through contaminant-name labels. A one-hot adsorbate label identifies the metal species but does not encode chemical similarity between ions. Descriptors such as charge, ionic radius, hydrated radius, ionic potential, and hydration ratio allow the model to partially capture differences in ion size, charge density, and hydration behavior that can influence adsorption.

The full model therefore used the following ML inputs: seven operating/material numeric variables; adsorption capacity, capacity-loading ratio, and thermo-capacity; five engineered process descriptors; six ion descriptors; normalized adsorbate identity; and broad adsorbent family. The capacity-free models retained the operating/material variables and engineered descriptors but removed adsorption capacity, capacity-loading ratio, and thermo-capacity. Among the capacity-free models, one retained ion descriptors, one used normalized adsorbate identity, and the strictest model used numeric variables only.

Capacity-derived predictors were included only in the capacity-inclusive process model and were excluded from all screening-safe models. They were defined as

$$\text{Capacity Loading Ratio} = \frac{q_e}{C_0}, \quad (9)$$

$$\text{Thermo Capacity} = q_e T, \quad (10)$$

where q_e is adsorption capacity, C_0 is initial concentration, and T is temperature.

2.4. Machine-learning models

XGBoost was selected as the primary model because gradient-boosted trees are effective for nonlinear tabular datasets and are widely used in adsorption ML studies [12, 13]. LightGBM was included as an additional gradient-boosting compara-

225 tor because preliminary model development indicated similar row-wise interpolation performance [14]. Extremely randomized trees (ExtraTrees) were included as a comparator ensemble model because random forest-style tree ensembles provide robust nonlinear baselines for tabular prediction [15, 16]. A mean dummy regressor was included to quantify improvement beyond a no-skill baseline.

230 Categorical variables were one-hot encoded with unknown-category handling, while numerical variables were passed directly to the tree models. For each split, categorical encoding and model fitting were performed within the training fold only, and the fitted preprocessing pipeline was then applied to the corresponding test fold. Unknown categories appearing only in test splits were ignored by the
235 encoder rather than fitted using test information. This ensured that preprocessing did not introduce information from the test set into the training process.

XGBoost used 140 estimators, learning rate 0.05, maximum depth 4, subsample 0.85, column subsample 0.85, ℓ_1 regularization 0.1, and ℓ_2 regularization 1.0. The standard LightGBM model used 300 estimators, learning rate 0.05, maximum
240 depth 5, 31 leaves, minimum child samples 15, subsample 0.85, column subsample 0.85, ℓ_1 regularization 0.1, and ℓ_2 regularization 0.5. ExtraTrees used 120 estimators, maximum depth 10, and minimum leaf size 2.

Two additional LightGBM sensitivity configurations were evaluated because preliminary model development had produced strong single-split LightGBM performance. These configurations were not used to claim a newly optimized algorithm;
245 they were included as sensitivity checks to test whether the validation conclusions were stable across reasonable gradient-boosting complexity settings. The conservative configuration used 600 estimators, learning rate 0.03, maximum depth 5, 20 leaves, minimum child samples 25, subsample 0.80, column subsample

250 0.80, ℓ_1 regularization 0.5, and ℓ_2 regularization 1.0. The expanded configuration used 500 estimators, learning rate 0.035, maximum depth 6, 40 leaves, minimum child samples 10, subsample 0.90, column subsample 0.90, ℓ_1 regularization 0.05, and ℓ_2 regularization 0.2. Predictions were clipped to the physically meaningful RE range of 0–100% before metric calculation. This clipping was applied only to
 255 final model predictions and did not affect training targets.

2.5. Validation design

Two validation tracks were used (Fig. 2). Track A evaluated source-specific interpolation. Although the cleaned dataset contained 189 normalized source references, most references contributed only a small number of rows. Therefore,
 260 source-specific 5-fold cross-validation was reported only for references with at least 20 usable records and sufficient variation in removal efficiency, using shuffled folds with a fixed random seed. Here, “5-fold” refers to the number of cross-validation folds within each eligible source group, not to the number of source groups reported. This tests whether high R^2 values are achievable within
 265 a known experimental source, which is similar to the validation setting of many high-performance adsorption ML studies.

Track B evaluated global generalization using three validation protocols within the restricted primary domain. The random row-wise split used an 80/20 train-test split stratified by the three retained pH-concentration regimes, so each split re-
 270 tained acidic/high- C_0 , acidic/low- C_0 , and neutral/high- C_0 observations. This split answers the in-dataset interpolation question: can the model predict held-out rows when related rows from the same sources may remain in training? This stratification affected only how rows were assigned to training and test sets; the regime label itself was not included in the fitted model. The adsorbent-held-out split used

group splitting so that exact adsorbent material names in the test set were absent from the training set; all rows for a held-out adsorbent identity were assigned to the test split. This split answers the cross-material question: can the model transfer to unseen material names when measured descriptors such as surface area, operating conditions, pH, dose, and concentration are available? It does not claim prediction for a completely undescribed material from name alone. The reference-held-out split used group splitting so that entire references/sources were absent from the training set, making it the stricter proxy for transfer across literature sources, experimental protocols, adsorbent synthesis routes, concentration windows, contaminant distributions, and reporting practices. This split answers the external-source question: can the model transfer to unseen literature/source domains? Because all data were literature-derived, reference-held-out validation was used as a proxy for external-source validation and should not be interpreted as a substitute for future prospective laboratory validation. Each Track B protocol was repeated over five random seeds. Metrics included R^2 , root mean squared error (RMSE), mean absolute error (MAE), and prediction bias.

3. Results

The results are reported in the same order as the validation framework: first, source-specific interpolation within coherent literature sources; second, repeated random row-wise prediction in the restricted primary domain; third, stricter adsorbent-held-out and reference-held-out transfer; and fourth, comparison of capacity-inclusive and capacity-free feature sets. This ordering separates interpolation accuracy from broader external-source generalization.

Source-aware validation framework for adsorption removal-efficiency prediction




Figure 2: Workflow of the proposed source-aware adsorption-ML validation framework. Track A measures source-specific interpolation, whereas Track B measures progressively stricter global generalization.

3.1. Source-specific models reproduced high adsorption-prediction accuracy

Track A showed that high removal-efficiency prediction accuracy is achievable when models are trained and evaluated within individual source groups. The source groups in Table 4 are not the only sources in the dataset; they are the clean bibliographic source groups that met the minimum row-count and target-variability criteria and produced stable positive within-source cross-validation performance. The full capacity process model achieved $R^2 = 0.9916$ for the source group Shen et al. (2017b), $R^2 = 0.9735$ for Gao et al. (2019), and $R^2 = 0.9601$ for Zama et al. (2017). These results are in the range commonly reported in adsorption ML studies under row-wise or within-domain validation [3, 5, 4]. They are also chemically plausible because the dominant source groups are internally coherent adsorption campaigns rather than arbitrary row collections: Shen et al. studied Ni(II) adsorption on wheat-straw-pellet and rice-husk biochars, Gao et al. investigated Cd(II) adsorption mechanisms on rice-straw and sewage-sludge biochars, Zama et al. studied Pb(II), Cd(II), and As(III) sorption/desorption across biochars, and Cui et al. focused on Cd(II) removal by *Canna indica*-derived biochar [17, 18, 19, 20]. Two additional ScienceDirect-indexed URL-bucket groups met the row-count threshold but were not treated as clean source-specific Track A groups because their reference metadata indicated mixed or unresolved bibliographic grouping. They are retained in the benchmark audit trail, but not used to support the high-interpolation claim. This reinforces the central point that source-specific predictability should be evaluated only for coherent source groups and should not be summarized by a single global accuracy value.

Capacity-free source-specific models also performed well in selected source groups. For example, the source-specific model for Zama et al. (2017) achieved

Table 4: Track A source-specific 5-fold cross-validation results for the four coherent source groups with stable positive within-source predictability. Complete Track A benchmark outputs, including unresolved ScienceDirect URL-bucket groups, are retained in the submitted CSV files for auditability.

Source group	Rows	R^2	RMSE	MAE
Shen et al. (2017b)	65	0.9916	3.20	2.33
Gao et al. (2019)	48	0.9735	2.97	2.14
Zama et al. (2017)	91	0.9601	7.84	3.56
Cui et al. (2016a)	24	0.8936	6.65	4.56




Figure 3: Track A source-specific interpolation performance. Several individual source-group models achieved R^2 values near or above 0.95, while other source groups remained more difficult.

$R^2 = 0.9250$ using the screening-safe ion feature set, while Shen et al. (2017b) achieved R^2 values above 0.91 using capacity-free screening feature sets. This indicates that high source-specific performance is not solely a capacity-inclusion artifact, although capacity-inclusive models generally gave the strongest interpolation performance.

3.2. Random row-wise validation produced moderate-to-good global performance

Under repeated random row-wise splits in the restricted primary domain, the full capacity process model achieved $R^2 = 0.8064$ and RMSE = 12.94 with LightGBM, while the corresponding XGBoost model achieved $R^2 = 0.7985$ and RMSE = 13.22. The screening-safe plus adsorbate one-hot XGBoost model achieved $R^2 = 0.7283$ and RMSE = 15.29, while the screening-safe plus ion XGBoost model achieved $R^2 = 0.7265$ and RMSE = 15.34. Thus, even without adsorption capacity, the model retained useful predictive ability under in-dataset interpolation (Table 5).

Table 5: Restricted-domain random row/regime split performance after excluding the sparse neutral/low- C_0 regime. Values are means over five repeated splits.

Feature/model	R^2	RMSE	MAE
Full capacity LightGBM	0.8064	12.94	8.08
Full capacity XGBoost	0.7985	13.22	8.44
Full capacity ExtraTrees	0.7730	14.02	9.02
Screening-safe plus adsorbate XGBoost	0.7283	15.29	10.84
Screening-safe plus ion XGBoost	0.7265	15.34	10.88

These results support the use of ML for in-domain adsorption prediction. They also explain the strong preliminary LightGBM results obtained during model development. Earlier LightGBM reports used a single random/regime-stratified split and included capacity-related predictors, giving test $R^2 = 0.847$ and RMSE = 12.05; in the present repeated-split benchmark, individual LightGBM random splits similarly reached $R^2 = 0.8499$ for the standard configuration and $R^2 = 0.8570$ for the conservative sensitivity configuration. However, the five-split means were lower ($R^2 = 0.8064$ and $R^2 = 0.7983$, respectively), showing that the earlier value was a valid favorable row-wise interpolation result rather than a stable estimate of source-held-out generalization. Random row-wise validation alone is therefore not sufficient evidence for deployment on unseen sources.

The two capacity-free XGBoost screening models differed in how the adsorbate was represented. The screening-safe plus adsorbate model used normalized adsorbate identity as a categorical one-hot feature, allowing the model to learn contaminant-specific effects directly. The screening-safe plus ion model used simple aqueous-ion descriptors, including charge, ionic radius, hydrated radius, ionic potential, and hydration ratio, instead of the direct contaminant label. The similar row-wise performances of these two models, $R^2 = 0.7283$ and $R^2 = 0.7265$, respectively, indicate that basic ion-chemistry descriptors captured much of the information provided by explicit adsorbate labels in the random-split setting.

3.3. Cross-material and cross-source validation tested broader transfer

The validation protocol strongly affected performance, which is expected for a heterogeneous dataset spanning many source references, adsorbent identities, and metal ions. In the adsorbent-held-out protocol, all experiments associated with selected adsorbent material names were removed from the training set and used

for testing. This is different from holding out adsorbates: the held-out groups are adsorbent materials, not contaminant ions. Under this cross-material test, the
365 best capacity-free result was obtained by the screening-safe numeric expanded-LightGBM model, with $R^2 = 0.5323$ and $\text{RMSE} = 19.35$. The screening-safe numeric XGBoost model achieved $R^2 = 0.5199$, while the conservative full-capacity LightGBM model achieved $R^2 = 0.5258$. These results indicate moderate but imperfect transfer to unseen adsorbent identities when measured descriptors and
370 operating conditions are available.

Reference-held-out validation was the most important broad-source stress test because every row from selected literature references was kept outside training. The model was therefore evaluated on source domains whose experimental protocol, adsorbent preparation, concentration window, contaminant distribution,
375 and reporting pattern were not represented by the same reference during training. This validation is closer to the practical question of whether a literature-trained model transfers to a new study. It was more difficult than row-wise validation but improved markedly after removing the sparse neutral/low- C_0 regime. The best reference-held-out sensitivity result was obtained by the full-capacity
380 expanded-LightGBM model, with $R^2 = 0.3296$ and $\text{RMSE} = 22.88$. Because this configuration includes capacity-related predictors, it should be interpreted as a process-monitoring or post-experiment result rather than as a pre-experiment screening result. The best capacity-free reference-held-out result remained the screening-safe numeric-only ExtraTrees model, with $R^2 = 0.3054$ and RMSE
385 $= 23.54$. These results remain below row-wise performance, but they show that defining a better-covered applicability domain produces more meaningful external-source validation for a broad, multi-source adsorption dataset.

Table 6: Restricted-domain Track B generalization performance after excluding the sparse neutral/low- C_0 regime. Values are means over five repeated splits.

Validation	Feature/model	R^2	RMSE	MAE
Holdout adsorbent	Screening-safe numeric expanded LightGBM	0.5323	19.35	14.55
Holdout adsorbent	Full capacity conservative LightGBM	0.5258	19.19	13.00
Holdout adsorbent	Screening-safe numeric-only XGBoost	0.5199	19.61	15.15
Holdout reference/source	Full capacity expanded LightGBM	0.3296	22.88	17.67
Holdout reference/source	Screening-safe numeric-only ExtraTrees	0.3054	23.54	18.67
Holdout reference/source	Full capacity standard LightGBM	0.2820	23.78	18.67




Figure 4: Track B validation comparison. Row-wise validation, adsorbent-held-out validation, and reference-held-out validation represent increasingly strict generalization tasks.




Figure 5: Validation cascade from source-specific interpolation to stricter global transfer. The figure summarizes why high R^2 values near 0.95 are achievable within known sources while reference-held-out prediction is a more difficult external-source task.

These results show that removal-efficiency prediction is strongly source-dependent. The same data and model family can produce high R^2 under source-specific interpolation, moderate-to-good R^2 under random global validation, and lower but still informative R^2 under source-held-out validation when the applicability domain is restricted to sufficiently represented pH-concentration regimes.

3.4. Source-domain heterogeneity explained validation sensitivity

The decrease from source-specific or row-wise performance to reference-held-out performance was not simply an algorithmic failure; it reflected the structure of the literature-compiled dataset. In the 537-record primary domain, the records were distributed across 180 normalized source references, but 142 of those sources contributed only one record and the median source size was one row (Table 7). Only six sources contributed at least 20 records, yet those six sources accounted for 52.7% of the primary-domain rows. Therefore, a random row-wise split can place closely related experiments from the same source campaign in both training and testing, whereas reference-held-out validation deliberately removes that source-specific support.

This audit identifies the hidden assumption behind the performance drop. Row-wise validation implicitly assumes that individual rows are exchangeable samples from one common adsorption distribution. The compiled dataset violates that assumption: source groups differ in contaminant identity, adsorbent synthesis route, pH and concentration window, measured material descriptors, analytical protocol, and reporting convention. The source-label variance diagnostic attributed 42.4% of the RE variance to source grouping, and the spread of source-wise mean RE values was larger than the median within-source spread. Thus, reference-held-out

Table 7: Source-domain heterogeneity diagnostics for the 537-record primary modeling domain.

Diagnostic	Value
Primary-domain records	537
Normalized source references	180
Singleton source references	142
Median records per source reference	1
Sources with at least 20 records	6
Rows contributed by six largest source groups	52.7%
Overall RE mean $\pm$ standard deviation	71.66 $\pm$ 29.92 percentage points
Standard deviation of source-wise RE means	19.26 percentage points
Median within-source RE standard deviation	7.37 percentage points
Source-label variance diagnostic, η^2_{source}	0.424

The source-label variance diagnostic is descriptive and was not used as a model input. Because many references were singletons, η^2_{source} should be interpreted as evidence of source imprint rather than as a causal partition of adsorption mechanisms.

validation is not merely a harsher score; it tests whether a model can transfer when source-specific experimental context is no longer shared between training and testing.
415

3.5. *Capacity-inclusive and capacity-free models answered different questions*

The full capacity process model performed best for random row-wise validation and several source-specific models. This is expected because adsorption capacity carries strong information about adsorption performance. However, adsorption capacity is derived after measuring concentration change and is usually not available before testing a candidate adsorbent under a new condition. Therefore, capacity-inclusive results should be interpreted as descriptive, process-monitoring, or post-experiment analysis results. Capacity-free models are more appropriate for pre-experiment screening and adsorbent-design guidance because they exclude adsorption capacity and capacity-derived predictors. In the restricted-domain adsorbent-held-out and reference-held-out tests, capacity-free numeric models were competitive or superior in several cases. This should not be interpreted as adsorption capacity being uninformative; rather, it suggests that capacity-related predictors can improve interpolation while also encoding source- or condition-specific relationships that do not always transfer cleanly across held-out materials or references.
420
425
430

4. Discussion

4.1. *Why near-0.95 R^2 is achievable but validation-dependent*

The Track A results explain why many adsorption ML studies report R^2 values near or above 0.95. Within a single source, adsorbent family, contaminant system, or experimental protocol, the response surface can be learned accurately by
435

tree-based models. In this setting, high R^2 is scientifically plausible and not automatically invalid. The present results reproduce that behavior, with source-specific R^2 values of 0.960–0.992 for three major source groups.

However, Track B shows that this performance should not be interpreted as
440 universal adsorption prediction. When sources are held out, the model must transfer across experimental protocols, adsorbent synthesis methods, contaminant distributions, concentration ranges, and reporting conventions. Under that validation setting, performance fell below row-wise performance even after excluding the sparse neutral/low- C_0 regime. The result is consistent with recent adsorption-
445 related work showing that random-split performance can be substantially higher than leave-compound-out performance [11].

The lower reference-held-out R^2 should therefore be read as a diagnostic finding rather than only a modeling weakness. It is not evidence that ML is unusable for adsorption; instead, it reveals the gap between interpolation within familiar
450 source domains and transfer to unseen literature/source domains. It also indicates that important context is missing from the current descriptor set. Likely missing variables include pH at point of zero charge, pore volume, pore-size distribution, surface functional groups, zeta potential, cation-exchange capacity, feedstock and pyrolysis details for biochars, aqueous speciation, ionic strength, competing
455 ions, and laboratory protocol metadata. Without such descriptors, the model can learn source-specific patterns under interpolation but has limited information for explaining why a new source behaves differently.

4.2. *Feature availability changes the scientific claim*

The comparison between capacity-inclusive and capacity-free feature sets
460 shows that feature availability changes the meaning of the prediction task.

Adsorption capacity carries strong information about removal efficiency because both are derived from concentration change during adsorption. Capacity-inclusive models are therefore useful for descriptive modeling, process monitoring, and post-experiment analysis, but they should not be interpreted as purely prospective
465 adsorbent-screening models unless capacity is already known or estimated independently.

In contrast, capacity-free models exclude adsorption capacity and capacity-derived terms, making them more appropriate for pre-experiment screening. These models rely on information that is typically available before testing, such as pH, initial
470 concentration, contact time, dose, temperature, surface area, adsorbent family, and adsorbate information. The similar row-wise performance of the adsorbate-label and ion-descriptor screening models indicates that simple ion-chemistry descriptors can capture much of the information provided by explicit contaminant labels in the random-split setting. The competitive performance of capacity-free
475 numeric models under held-out validation further suggests that simpler descriptor sets may sometimes transfer more robustly across materials or sources, even if capacity-inclusive models remain stronger for within-domain interpolation.

4.3. Experimental design needed to improve capacity-free screening

If the intended use is pre-experiment adsorbent screening, capacity-free models
480 are the more defensible modeling pathway because they use only information that can be known before measuring the final concentration change. The lower held-out R^2 values in the present study therefore indicate what future experiments must measure more systematically. A capacity-free model can improve only if the variables that control adsorption are available as inputs before the adsorption test,
485 rather than being indirectly supplied through post-experiment capacity.

Future datasets should therefore be designed as prospective ML benchmarks rather than assembled only from optimized adsorption case studies. Each adsorbent should be characterized before adsorption using descriptors such as BET surface area, pore volume, pore-size distribution, pH at point of zero charge, zeta potential, cation-exchange capacity, elemental composition, surface functional groups from FTIR/XPS, feedstock, pyrolysis or activation temperature, and modification chemistry. Solution chemistry should be recorded with similar detail, including initial pH, initial contaminant concentration, sulfate concentration, ionic strength, competing ions, alkalinity, redox condition where relevant, and whether the water is synthetic single-ion solution, synthetic AMD, or real AMD. These descriptors would give a capacity-free model direct chemical information that is currently hidden inside source-specific patterns.

The experimental matrix should also be balanced across the operating space. Instead of reporting mainly high-performing or optimized conditions, future studies should deliberately sample low, medium, and high initial concentrations; acidic and near-neutral pH; multiple adsorbent doses; several contact times; and multiple metal ions or mixed-metal AMD-relevant systems. Poorly represented regions, such as the neutral/low- C_0 regime in this study, should be intentionally filled rather than excluded after modeling. Replicate experiments should be included to estimate experimental noise, and a subset of adsorbents, contaminants, and laboratory batches should be reserved as blinded holdout groups before model fitting. This design would reduce random-split optimism and make adsorbent-held-out, contaminant-held-out, laboratory-held-out, and reference-held-out validation more stable.

In this workflow, adsorption capacity should still be measured and reported,

but it should be treated as an outcome or diagnostic variable rather than a predictor for screening. A practical future protocol would first train a capacity-free model using only pre-test descriptors to propose promising adsorbent–condition combinations, then experimentally measure both RE and capacity, and finally use the new results to update the model through active learning. Such an iterative design would improve R^2 not by hiding outcome-derived information in the feature set, but by expanding the chemically relevant input space and reducing extrapolation to poorly described sources.

4.4. Relation to existing adsorption ML literature

The proposed framework complements rather than contradicts previous high-accuracy adsorption ML studies. Hafsa et al. demonstrated that tree-based algorithms can achieve high removal-efficiency prediction accuracy for heavy-metal adsorption datasets under repeated cross-validation [3]. Wang et al. showed that XGBoost, SHAP, and partial-dependence analyses can identify interpretable drivers of heavy-metal removal by biochar [5]. Lu et al. showed that random forests can predict heavy-metal removal by chitosan-based flocculants with R^2 above 0.93 [4]. Barkhordari et al. further demonstrated that interpretable XGBoost models are relevant to JECE’s environmental-remediation scope by predicting heavy-metal removal efficiency in electrokinetic remediation [7]. These studies establish that ML is useful for environmental treatment modeling.

The present study adds a validation-layer contribution: adsorption ML workflows should explicitly state whether they are predicting within a known source, interpolating within a compiled dataset, transferring to unseen adsorbent names, or transferring to unseen sources. This distinction is critical because high row-wise R^2 can be overoptimistic when the deployment target is a new laboratory, adsor-

bent synthesis route, contaminant matrix, or literature source. The broader ML literature has identified leakage and validation mismatch as major reproducibility risks [9, 10], and adsorption datasets have several mechanisms by which such mismatch can occur.

540 4.5. *Practical implications for adsorption model deployment*

For practical adsorption screening, a single universal RE model should be used cautiously. A more defensible workflow is to define the pH-concentration applicability domain before modeling, exclude or separately model poorly represented regimes, use source-specific or laboratory-calibrated models for high-accuracy optimization within a known experimental system, use capacity-free screening models for preliminary adsorbent and operating-condition selection, and report source-held-out validation when predictions are intended for new adsorbents, contaminants, laboratories, or literature domains. This framing allows high- R^2 models to remain useful without overstating their generality.

550 The most direct route to improving cross-source performance is not simply replacing XGBoost or LightGBM with a more complex learner. The source audit suggests that external performance would more likely improve through richer mechanistic descriptors, better source metadata, and validation-aware modeling. Future versions of the dataset should prioritize adsorbent surface chemistry, synthesis and activation conditions, solution chemistry, and explicit laboratory/source labels. Once those descriptors are available, hierarchical or multi-task models could separate universal adsorption trends from source-specific offsets, while still reporting reference-held-out or laboratory-held-out validation as the primary deployment test.

560 4.6. *Limitations and future work*

This study has several limitations. First, the dataset represents adsorption tests for AMD-relevant metal and metalloid contaminants, but many records are single- or limited-ion aqueous experiments rather than complete real AMD matrices. Therefore, sulfate concentration, ionic strength, competing ions, 565 dissolved organic matter, and site-specific mine-water composition should be added before claiming direct deployment to complex field AMD. Second, the dataset does not contain explicit laboratory-origin labels for all records; bibliographic reference was therefore used as a source proxy. Future work should annotate records as in-house, external-laboratory, or literature-derived and repeat 570 validation with leave-laboratory-out splits. Third, the ion descriptor table should be expanded with fully cited descriptors such as electronegativity, hydration energy, softness/hardness, and likely aqueous speciation. Fourth, important adsorbent descriptors such as pH_{pzc} , pore volume, pore size, zeta potential, FTIR-derived functional groups, and cation-exchange capacity are incomplete in 575 the present dataset. Fifth, the neutral/low-concentration regime contained only 26 observations and was excluded from primary modeling; predictions in that regime should therefore be treated as outside the present applicability domain until more data are available. Sixth, because all records were literature-derived, reference-held-out validation should be interpreted as a proxy for external-source transfer, 580 not as a substitute for prospective laboratory validation. Finally, the dataset was designed as a heterogeneous ML benchmark for source-aware validation rather than as a complete systematic meta-analysis of all AMD treatment or heavy-metal adsorption studies. These limitations define clear routes for improving external predictive performance while preserving the source-aware validation principle.

585 5. Conclusions

This study developed a source-aware and leakage-aware validation framework plus applicability-domain audit for AMD-relevant metal-contaminant adsorption removal-efficiency prediction. It should not be read as a claim of a state-of-the-art universal adsorption predictor. High R^2 values were achievable within
590 individual adsorption sources, with source-specific XGBoost models reaching R^2 up to 0.992. Because the neutral/low-concentration regime contained only 26 observations, primary global modeling excluded that sparse regime and used a 537-record restricted domain. Within this domain, random row-wise validation reached $R^2 = 0.806$ for the full capacity LightGBM process model, adsorbent-held-out validation reached $R^2 = 0.532$ for the capacity-free numeric
595 expanded-LightGBM screening model, and reference-held-out validation reached $R^2 = 0.330$ for the full-capacity expanded-LightGBM sensitivity model. The best capacity-free reference-held-out model reached $R^2 = 0.305$. These findings show that high-accuracy interpolation, applicability-domain definition, and
600 cross-source generalization are different scientific tasks. High adsorption-ML accuracy is meaningful when the claim is interpolation within known sources or well-represented operating domains, but it is not sufficient by itself to claim deployment to unseen adsorbents, unseen studies, or new laboratories. Future adsorption ML studies should explicitly state the validation question being
605 answered, separate capacity-inclusive process-monitoring models from capacity-free screening models, and include source- or laboratory-held-out validation when deployment beyond the training source is claimed.

CRedit authorship contribution statement

Yash Bisht: Conceptualization, Methodology, Software, Formal analysis, Vi-
610 sualization, Writing – original draft. Margreth Tadie: Supervision, Validation,
Writing – review and editing. Susmit Chitransh: Data curation, Conceptualiza-
tion, Investigation, Validation, Writing – review and editing.

Declaration of competing interest

The authors declare that they have no known competing financial interests or
615 personal relationships that could have appeared to influence the work reported in
this paper.

Funding

This research did not receive any specific grant from funding agencies in the
public, commercial, or not-for-profit sectors.

620 Data availability

The curated dataset, analysis scripts, and generated benchmark outputs are
available from the authors upon reasonable request and will be deposited in a pub-
lic repository before final publication.

625 Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During the preparation of this work, the authors used OpenAI’s Codex to as-
sist with code organization, LaTeX formatting, and language polishing. After us-

ing this tool, the authors reviewed and edited the content as needed and take full responsibility for the content of the manuscript.

630 **Acknowledgements**

The authors acknowledge the institutional support of the Indian Institute of Technology Roorkee and Stellenbosch University.

References

- [1] G.S. Simate, S. Ndlovu, Acid mine drainage: Challenges and opportunities, *Journal of Environmental Chemical Engineering* 2 (2014) 1785–1803.
635 <https://doi.org/10.1016/j.jece.2014.07.021>.
- [2] V. Dube, Z. Phiri, A.T. Kuvarega, B.B. Mamba, L.-A. de Kock, Exploring acid mine drainage treatment through adsorption: a bibliometric analysis, *Environmental Science and Pollution Research* 31 (2024) 59659–59680.
640 <https://doi.org/10.1007/s11356-024-35047-2>.
- [3] N. Hafsa, S. Rushd, M. Al-Yaari, M. Rahman, A generalized method for modeling the adsorption of heavy metals with machine learning algorithms, *Water* 12 (2020) 3490. <https://doi.org/10.3390/w12123490>.
- [4] C. Lu, Z. Xu, B. Dong, Y. Zhang, M. Wang, Y. Zeng, C. Zhang, Machine
645 learning for the prediction of heavy metal removal by chitosan-based flocculants, *Carbohydrate Polymers* 285 (2022) 119240. <https://doi.org/10.1016/j.carbpol.2022.119240>.

- [5] C. Wang, Y. Zhao, Y. Gao, H. Chen, X. Li, B. Zhou, D. Fan, Z. Fang, J. Liu, Interpretable machine learning for predicting heavy metal removal and optimizing biochar characteristics, *Journal of Water Process Engineering* 68 (2024) 106484. <https://doi.org/10.1016/j.jwpe.2024.106484>.
650
- [6] Z.M. Yaseen, F.L. Alhalimi, Heavy metal adsorption efficiency prediction using biochar properties: a comparative analysis for ensemble machine learning models, *Scientific Reports* 15 (2025). <https://doi.org/10.1038/s41598-025-96271-5>.
655
- [7] M.S. Barkhordari, N. Zhou, K. Li, C. Qi, Interpretable machine learning for predicting heavy metal removal efficiency in electrokinetic soil remediation, *Journal of Environmental Chemical Engineering* 12 (2024) 114330. <https://doi.org/10.1016/j.jece.2024.114330>.
- [8] S.M. Lundberg, S.-I. Lee, A unified approach to interpreting model predictions, in: *Advances in Neural Information Processing Systems* 30, 2017, pp. 4765–4774.
660
- [9] S. Kapoor, A. Narayanan, Leakage and the reproducibility crisis in machine-learning-based science, *Patterns* 4 (2023) 100804. <https://doi.org/10.1016/j.patter.2023.100804>.
665
- [10] J.-J. Zhu, M. Yang, Z.J. Ren, Machine learning in environmental research: Common pitfalls and best practices, *Environmental Science & Technology* 57 (2023) 17671–17689. <https://doi.org/10.1021/acs.est.3c00026>.
- [11] H.V. Patel, J. Green, H. Park, S. Luster-Teasley Pass, R. Zhao, A hybrid response surface methodology and machine learning framework for quan-
670

tifying effects of physicochemical parameters on PFAS distribution, ACS ES&T Water (2026) Advance online publication. <https://doi.org/10.1021/acsestwater.5c01162>.

- [12] T. Chen, C. Guestrin, XGBoost: A scalable tree boosting system, in: Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, ACM, San Francisco, CA, USA, 2016, pp. 785–794. <https://doi.org/10.1145/2939672.2939785>.
675
- [13] J.H. Friedman, Greedy function approximation: A gradient boosting machine, Annals of Statistics 29 (2001) 1189–1232. <https://doi.org/10.1214/aos/1013203451>.
680
- [14] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, T.-Y. Liu, LightGBM: A highly efficient gradient boosting decision tree, in: Advances in Neural Information Processing Systems 30, 2017, pp. 3146–3154.
- [15] L. Breiman, Random forests, Machine Learning 45 (2001) 5–32. <https://doi.org/10.1023/A:1010933404324>.
685
- [16] P. Geurts, D. Ernst, L. Wehenkel, Extremely randomized trees, Machine Learning 63 (2006) 3–42. <https://doi.org/10.1007/s10994-006-6226-1>.
- [17] Z. Shen, Y. Zhang, O. McMillan, F. Jin, A. Al-Tabbaa, Characteristics and mechanisms of nickel adsorption on biochars produced from wheat straw pellets and rice husk, Environmental Science and Pollution Research 24 (2017) 12809–12819. <https://doi.org/10.1007/s11356-017-8847-2>.
690

- [18] L.-Y. Gao, J.-H. Deng, G.-F. Huang, K. Li, K.-Z. Cai, Y. Liu, F. Huang, Relative distribution of Cd(II) adsorption mechanisms on biochars derived from rice straw and sewage sludge, *Bioresource Technology* 272 (2019) 114–122. <https://doi.org/10.1016/j.biortech.2018.09.138>.
695
- [19] E.F. Zama, Y.-G. Zhu, B.J. Reid, G.-X. Sun, The role of biochar properties in influencing the sorption and desorption of Pb(II), Cd(II) and As(III) in aqueous solution, *Journal of Cleaner Production* 148 (2017) 127–136. <https://doi.org/10.1016/j.jclepro.2017.01.125>.
700
- [20] X. Cui, S. Fang, Y. Yao, T. Li, Q. Ni, X. Yang, Z. He, Potential mechanisms of cadmium removal from aqueous solution by *Canna indica* derived biochar, *Science of the Total Environment* 562 (2016) 517–525. <https://doi.org/10.1016/j.scitotenv.2016.03.248>.