


Graphical Abstract

Adsorption ML needs source-aware validation



High interpolation accuracy, applicability domain, and external generalization are different scientific claims.

Highlights

- Source-specific adsorption models reached R^2 up to 0.992.
- Sparse neutral-low-concentration cases were excluded from primary modeling.
- Row-wise validation overstated source-transfer performance.
- Reference-held-out screening reached $R^2 = 0.305$ in the main domain.
- Capacity-inclusive and capacity-free models served distinct uses.

Source-aware machine learning for heavy-metal adsorption: distinguishing high-accuracy interpolation from cross-source generalization

Yash Bisht^{a,*}, Susmit Chitransh^b

^a*Department of Mechanical and Industrial Engineering, Indian Institute of Technology Roorkee, Roorkee 247667, Uttarakhand, India*

^b*Department of Chemical Engineering, University of Cape Town, Rondebosch 7700, Cape Town, South Africa*

Abstract

Machine learning is increasingly used to predict adsorption removal efficiency for water and wastewater treatment, but reported performance can depend strongly on feature definition, domain coverage, and validation design. This study develops a source-aware and leakage-aware framework for heavy-metal adsorption removal-efficiency prediction using a literature-compiled experimental dataset. The raw dataset contained 566 adsorption records; after deterministic numeric parsing and quality filtering, 563 valid records remained. Because the neutral/low-concentration regime contained only 26 observations, primary modeling was restricted to the remaining 537 records covering acidic/high-concentration, acidic/low-concentration, and neutral/high-concentration regimes. Four feature sets were evaluated to separate capacity-inclusive process-monitoring models from capacity-free pre-experiment screening models. Track A quantified source-specific interpolation through within-source cross-validation, where

*Corresponding author.

XGBoost reached R^2 values of 0.992, 0.974, and 0.960 for three major source groups. Track B evaluated broader generalization through repeated random row-wise, adsorbent-held-out, and reference-held-out splits. In the restricted primary domain, the full capacity-inclusive LightGBM model achieved $R^2 = 0.806$ under random row-wise validation. The best adsorbent-held-out result was obtained by a capacity-free numeric XGBoost model ($R^2 = 0.520$), and the best reference-held-out result was obtained by a capacity-free numeric ExtraTrees model ($R^2 = 0.305$). These results show that high adsorption-prediction accuracy is meaningful for in-source interpolation, but broader deployment requires explicit applicability-domain definition and source-held-out validation.

Keywords: adsorption, heavy metals, removal efficiency, machine learning, XGBoost, source-aware validation, wastewater treatment

1. Introduction

Heavy-metal contamination remains a persistent challenge for water and wastewater treatment because metals such as Pb, Cd, Cr, Ni, Cu, As, Hg, and Zn are toxic, non-biodegradable, and prone to accumulation in environmental and biological systems. Adsorption is widely investigated for heavy-metal removal because it can operate under mild conditions, can be adapted to low-cost or waste-derived adsorbents, and can be optimized through pH, adsorbent dose, contact time, temperature, surface chemistry, and initial contaminant concentration. The expansion of experimental adsorption datasets has therefore encouraged the use of machine learning (ML) to predict removal efficiency and accelerate adsorbent screening.

Recent adsorption and environmental-remediation ML studies demonstrate

that high predictive accuracy is possible. Hafsa et al. modeled heavy-metal adsorption efficiencies using multiple ML algorithms and reported R^2 values in the 0.96–0.99 range under ten-fold cross-validation, with experimental data supplemented by synthetic interpolation [1]. Lu et al. predicted heavy-metal removal by chitosan-based flocculants using random forests and reported $R^2 = 0.9354$, identifying solution pH and flocculant molecular weight as important drivers [2]. Wang et al. used interpretable ML for heavy-metal removal by biochar and reported XGBoost test performance near $R^2 = 0.99$, together with SHAP and partial-dependence analyses to relate removal efficiency to biochar and operating variables [3, 6]. Recent ensemble-learning work on biochar adsorption similarly reported strong XGBoost performance for heavy-metal adsorption efficiency prediction [4]. Within the scope of *Journal of Environmental Chemical Engineering*, interpretable XGBoost modeling has also been applied to heavy-metal removal efficiency in electrokinetic remediation, demonstrating the journal relevance of transparent ML workflows for environmental treatment systems [5].

The central challenge is that high row-wise validation performance does not necessarily imply generalization to unseen laboratories, adsorbents, pollutant systems, or source studies. In ML-based science, data leakage and validation mismatch can produce overoptimistic estimates of deployable performance [7]. Environmental ML best-practice guidance similarly emphasizes careful data splitting, leakage management, randomness assessment, and evaluation design [8]. These concerns are particularly important for adsorption datasets because different sources may use different adsorbent synthesis routes, contaminant species, concentration ranges, analytical protocols, and reporting conventions. A random row-wise split can place related observations from the same reference, adsorbent

family, or operating campaign into both the training and test sets. Such a model may perform well, but the result primarily measures interpolation within a familiar
 40 experimental domain.

Feature definition can further blur the interpretation of reported accuracy. Adsorption capacity and removal efficiency are both derived from concentration changes in many batch adsorption experiments:

$$\text{RE}(\%) = 100 \frac{C_0 - C_e}{C_0}, \quad (1)$$

$$q_e = \frac{(C_0 - C_e)V}{m}, \quad (2)$$

where C_0 and C_e are the initial and equilibrium concentrations, V is the solu-
 45 tion volume, and m is the adsorbent mass. Capacity can therefore be useful for process monitoring after adsorption has been measured, but it should be treated cautiously in pre-experiment screening models where capacity is not yet available. Recent PFAS adsorption modeling has similarly shown that random-split performance can be much stronger than leave-one-compound-out performance, illustrating that interpolation and transfer to unseen chemical domains are distinct tasks
 50 [9].

This study addresses these issues by developing a dual-track validation framework for heavy-metal adsorption removal-efficiency prediction across a broad literature-derived dataset spanning 189 normalized source references,
 55 219 adsorbents, and 16 adsorbate classes. Track A evaluates source-specific interpolation and asks whether high R^2 values comparable to the adsorption-ML literature can be reproduced within individual source groups. Track B evaluates global generalization under random row-wise, adsorbent-held-out, and

reference-held-out validation. The study further separates capacity-inclusive
60 process-monitoring models from capacity-free screening models. The novelty
is not simply a new regression model; rather, it is a source-aware reporting
framework that clarifies when high R^2 is scientifically valid, when it is validation-
dependent, and how adsorption ML studies should present performance for
practical environmental-engineering use.

65 **2. Materials and methods**

2.1. Dataset and target variable

The dataset was compiled from reported heavy-metal adsorption experiments.
Each record corresponds to one adsorption condition with a reported removal ef-
ficiency. The raw dataset contained 566 rows and 13 columns. After numeric
70 parsing, missing-value filtering, and constraining removal efficiency to the phys-
ically meaningful range 0–100%, 563 records remained. The cleaned dataset in-
cluded 189 normalized source references, 219 adsorbents, and 16 normalized ad-
sorbate classes. The primary modeling domain excluded the sparse neutral/low-
concentration regime, leaving 537 records for model training and testing.

75 The target variable was removal efficiency, RE (%). Input variables included
adsorbent name, adsorbate identity, surface area, adsorption capacity, initial con-
centration, contact time, adsorbent dose, agitation speed, initial pH, temperature,
and source reference. The source identifier, denoted Ref_group, was used as a
source-origin proxy because explicit laboratory-origin labels were not available
80 for all records.

2.2. *Numeric cleaning and categorical normalization*

Literature-derived adsorption tables often contain nonuniform numeric formatting. Numeric entries were cleaned using deterministic rules: decimal commas were converted to decimal points, uncertainty and approximate symbols were removed, multiple values in one cell were reduced to the first numeric value, and nonnumeric text was parsed using regular expressions. Rows with missing values in required numeric variables were removed. Adsorbate names were normalized to reduce duplicates such as elemental and valence-state aliases. Adsorbent names were grouped into broad families, including carbon-based, mineral/oxide, polymer, biomass, and composite/other categories.

The cleaned data were partitioned into four pH-concentration regimes using $\text{pH} = 7$ as the acidic/near-neutral boundary and the dataset median initial concentration, $C_0 = 50$ ppm, as the low/high concentration boundary (Table 1). The median C_0 is the middle value of the initial-concentration distribution and provides a data-driven threshold for separating lower- and higher-concentration experiments. These regimes were not used as prediction targets. Instead, they were used for stratified random splitting and diagnostic analysis, ensuring that train-test splits preserved chemically meaningful coverage across major operating regions. The acidic regimes dominated the dataset, whereas the neutral/low- C_0 regime contained only 26 observations. Because this regime was too small for stable split-wise evaluation, it was excluded from the primary modeling domain and treated as outside the present applicability domain.

2.3. *Feature sets*

Four feature sets were constructed to distinguish scientific use cases (Table 2). This separation is important because a model that includes adsorption capacity

Table 1: pH-concentration regime distribution in the cleaned dataset.

Regime	Count	Mean pH	Mean C_0 (ppm)	Mean RE (%)
Acidic_HighC0	246	4.51	190.27	70.30
Acidic_LowC0	231	5.12	11.04	74.06
Neutral_HighC0	60	7.65	166.51	68.05
Neutral_LowC0	26	7.47	20.24	81.41

Neutral_LowC0 was excluded from primary modeling because it contained only 26 observations; the other three regimes formed the 537-record primary domain.

answers a different question from a model intended for pre-experiment adsorbent screening. The base experimental inputs were surface area, initial concentration, contact time, dose, agitation speed, initial pH, and temperature. The process-monitoring feature set additionally used adsorption capacity and capacity-derived terms, whereas the screening feature sets deliberately removed these terms.

Physics-inspired variables included site density, logarithmic contact time, mixing index, driving-force proxy, and acidity strength:

$$\text{Site Density} = \frac{\text{dose} \times \text{surface area}}{C_0}, \quad (3)$$

$$\text{LogTime} = \log(1 + \text{contact time}), \quad (4)$$

$$\text{Mixing Index} = \text{RPM} \times \text{LogTime}, \quad (5)$$

$$\text{Driving Force} = \frac{C_0}{\text{dose}}, \quad (6)$$

$$\text{Acidity Strength} = |\text{pH} - 7|. \quad (7)$$

Ion descriptors were added to represent adsorbates through simple aqueous-chemistry properties rather than only through pollutant-name labels. A one-hot

Table 2: Feature-set definitions used to separate process-monitoring and screening tasks.

Feature set	Included information	Intended use
Full capacity process model	Operating variables, surface area, adsorption capacity, capacity-derived ratios, ion descriptors, adsorbent family, and adsorbate identity	Descriptive or process-monitoring prediction
Screening-safe plus ion	Operating variables, surface area, engineered process variables, ion descriptors, and adsorbent family; excludes capacity terms	Pre-experiment screening with ion chemistry
Screening-safe plus adsorbate one-hot	Operating variables, surface area, engineered process variables, adsorbate identity, and adsorbent family; excludes capacity terms	Screening where pollutant identity is known
Screening-safe numeric-only	Operating variables, surface area, and engineered process variables; excludes capacity and adsorbate identity	Strict screening baseline

115 adsorbate label identifies the metal species but does not encode chemical similarity between ions. Descriptors such as charge, ionic radius, hydrated radius, ionic potential, and hydration ratio allow the model to partially capture differences in ion size, charge density, and hydration behavior that can influence adsorption.

The full model therefore used the following ML inputs: seven operating/material numeric variables; adsorption capacity, capacity-loading ratio, and thermo-capacity; five engineered process descriptors; six ion descriptors; normalized adsorbate identity; and broad adsorbent family. The capacity-free models retained the operating/material variables and engineered descriptors but removed adsorption capacity, capacity-loading ratio, and thermo-capacity. Among the capacity-free models, one retained ion descriptors, one used normalized adsorbate identity, and the strictest model used numeric variables only.

2.4. Machine-learning models

XGBoost was selected as the primary model because gradient-boosted trees are effective for nonlinear tabular datasets and are widely used in adsorption ML studies [10, 11]. LightGBM was included as an additional gradient-boosting comparator because preliminary model development indicated similar row-wise interpolation performance. Extremely randomized trees (ExtraTrees) were included as a comparator ensemble model because random forest-style tree ensembles provide robust nonlinear baselines for tabular prediction [12, 13]. A mean dummy regressor was included to quantify improvement beyond a no-skill baseline.

Categorical variables were one-hot encoded with unknown-category handling, while numerical variables were passed directly to the tree models. XGBoost used 140 estimators, learning rate 0.05, maximum depth 4, subsample 0.85, column subsample 0.85, ℓ_1 regularization 0.1, and ℓ_2 regularization 1.0. LightGBM used

140 300 estimators, learning rate 0.05, maximum depth 5, 31 leaves, minimum child
samples 15, subsample 0.85, column subsample 0.85, ℓ_1 regularization 0.1, and
 ℓ_2 regularization 0.5. ExtraTrees used 120 estimators, maximum depth 10, and
minimum leaf size 2. Predictions were clipped to the physically meaningful RE
range of 0–100% before metric calculation.

145 2.5. Validation design

Two validation tracks were used (Fig. 1). Track A evaluated source-specific
interpolation. Although the cleaned dataset contained 189 normalized source ref-
erences, most references contributed only a small number of rows. Therefore,
source-specific 5-fold cross-validation was reported only for references with at
150 least 20 usable records and sufficient variation in removal efficiency, using shuf-
fled folds with a fixed random seed. This tests whether high R^2 values are achiev-
able within a known experimental source, which is similar to the validation setting
of many high-performance adsorption ML studies.

Track B evaluated global generalization using three validation protocols
155 within the restricted primary domain. The random row-wise split used an 80/20
train-test split stratified by the three retained pH-concentration regimes, so each
split retained acidic/high- C_0 , acidic/low- C_0 , and neutral/high- C_0 observations.
The adsorbent-held-out split used group splitting so that exact adsorbent material
names in the test set were absent from the training set; all rows for a held-out
160 adsorbent identity were assigned to the test split. This protocol tests transfer to un-
seen material identities when measured descriptors such as surface area, operating
conditions, pH, dose, and concentration are available. It does not claim prediction
for a completely undescribed material from name alone. The reference-held-out
split used group splitting so that entire references/sources were absent from the

165 training set, making it the stricter proxy for transfer across literature sources, experimental protocols, adsorbent synthesis routes, and reporting practices. Each Track B protocol was repeated over five random seeds. Metrics included R^2 , root mean squared error (RMSE), mean absolute error (MAE), and prediction bias.

Source-aware validation framework for adsorption removal-efficiency prediction




Figure 1: Workflow of the proposed source-aware adsorption-ML validation framework. Track A measures source-specific interpolation, whereas Track B measures progressively stricter global generalization.

3. Results

170 3.1. Source-specific models reproduced high adsorption-prediction accuracy

Track A showed that high removal-efficiency prediction accuracy is achievable when models are trained and evaluated within individual source groups. The source groups in Table 3 are not the only sources in the dataset; they are the source

groups that met the minimum row-count and target-variability criteria for within-
 175 source cross-validation. The full capacity process model achieved $R^2 = 0.9916$
 for the source group Shen et al. (2017b), $R^2 = 0.9735$ for Gao et al. (2019), and
 $R^2 = 0.9601$ for Zama et al. (2017). These results are in the range commonly
 reported in adsorption ML studies under row-wise or within-domain validation
 [1, 3, 2]. The lower-performing eligible source groups show that source-specific
 180 predictability is not uniform, which further supports reporting source-aware vali-
 dation rather than a single global accuracy value.

Table 3: Track A source-specific 5-fold cross-validation results for reported source groups that satisfied the minimum row-count and target-variability criteria.

Source group	Rows	R^2	RMSE	MAE
Shen et al. (2017b)	65	0.9916	3.20	2.33
Gao et al. (2019)	48	0.9735	2.97	2.14
Zama et al. (2017)	91	0.9601	7.84	3.56
Cui et al. (2016a)	24	0.8936	6.65	4.56
ScienceDirect source PII:S2213343721006655	30	-0.5584	10.45	6.88
ScienceDirect source PII:S0045653519308616	25	-1.6159	26.70	18.86

Capacity-free source-specific models also performed well in selected source
 groups. For example, the source-specific model for Zama et al. (2017) achieved
 $R^2 = 0.9250$ using the screening-safe ion feature set, while Shen et al. (2017b)
 185 achieved R^2 values above 0.91 using capacity-free screening feature sets. This
 indicates that high source-specific performance is not solely a capacity-inclusion




Figure 2: Track A source-specific interpolation performance. Several individual source-group models achieved R^2 values near or above 0.95, while other source groups remained more difficult.

artifact, although capacity-inclusive models generally gave the strongest interpolation performance.

3.2. Random row-wise validation produced moderate-to-good global performance

190

Under repeated random row-wise splits in the restricted primary domain, the full capacity process model achieved $R^2 = 0.8064$ and RMSE = 12.94 with LightGBM, while the corresponding XGBoost model achieved $R^2 = 0.7985$ and RMSE = 13.22. The screening-safe plus adsorbate one-hot XGBoost model achieved $R^2 = 0.7283$ and RMSE = 15.29, while the screening-safe plus ion XGBoost model achieved $R^2 = 0.7265$ and RMSE = 15.34. Thus, even without adsorption capacity, the model retained useful predictive ability under in-dataset interpolation (Table 4).

Table 4: Restricted-domain random row/regime split performance after excluding the sparse neutral/low- C_0 regime. Values are means over five repeated splits.

Feature/model	R^2	RMSE	MAE
Full capacity LightGBM	0.8064	12.94	8.08
Full capacity XGBoost	0.7985	13.22	8.44
Full capacity ExtraTrees	0.7730	14.02	9.02
Screening-safe plus adsorbate XGBoost	0.7283	15.29	10.84
Screening-safe plus ion XGBoost	0.7265	15.34	10.88

These results support the use of ML for in-domain adsorption prediction. They also explain the strong preliminary LightGBM results obtained during model development: earlier LightGBM runs used the same type of random/regime-stratified split and included capacity-related predictors, giving global R^2 values around 0.84–0.86 depending on the exact feature set and domain filter. Those values are valid row-wise interpolation results, but they are not evidence of source-held-out generalization. Random row-wise validation alone is therefore not sufficient evidence for deployment on unseen sources.

The two capacity-free XGBoost screening models differed in how the adsorbate was represented. The screening-safe plus adsorbate model used normalized adsorbate identity as a categorical one-hot feature, allowing the model to learn pollutant-specific effects directly. The screening-safe plus ion model used simple aqueous-ion descriptors, including charge, ionic radius, hydrated radius, ionic potential, and hydration ratio, instead of the direct pollutant label. The similar row-wise performances of these two models, $R^2 = 0.7283$ and $R^2 = 0.7265$, respectively, indicate that basic ion-chemistry descriptors captured much of the

215 information provided by explicit adsorbate labels in the random-split setting.

3.3. *Cross-material and cross-source validation tested broader transfer*

The validation protocol strongly affected performance, which is expected for a heterogeneous dataset spanning many source references, adsorbent identities, and metal ions. In the adsorbent-held-out protocol, all experiments associated with
220 selected adsorbent material names were removed from the training set and used for testing. This is different from holding out adsorbates: the held-out groups are adsorbent materials, not pollutant ions. Under this cross-material test, the best result was obtained by the screening-safe numeric-only XGBoost model, with $R^2 = 0.5199$ and $RMSE = 19.61$. The corresponding LightGBM screening model
225 achieved $R^2 = 0.5002$, while the full capacity XGBoost process model achieved $R^2 = 0.4385$. These results indicate moderate but imperfect transfer to unseen adsorbent identities when measured descriptors and operating conditions are available.

Reference-held-out validation was the most important broad-source stress test
230 because every row from selected literature references was kept outside training. The model was therefore evaluated on source domains whose experimental protocol, adsorbent preparation, concentration window, pollutant distribution, and reporting pattern were not represented by the same reference during training. This validation is closer to the practical question of whether a literature-trained
235 model transfers to a new study. It was more difficult than row-wise validation but improved markedly after removing the sparse neutral/low- C_0 regime. The best reference-held-out result was obtained by the screening-safe numeric-only Extra-Trees model, with $R^2 = 0.3054$ and $RMSE = 23.54$. The full capacity LightGBM model achieved $R^2 = 0.2820$, while the full capacity XGBoost model achieved

240 $R^2 = 0.1678$. These results remain below row-wise performance, but they show that defining a better-covered applicability domain produces more meaningful external-source validation for a broad, multi-source adsorption dataset.

Table 5: Restricted-domain Track B generalization performance after excluding the sparse neutral/low- C_0 regime. Values are means over five repeated splits.

Validation	Feature/model	R^2	RMSE	MAE
Holdout adsorbent	Screening-safe numeric-only XGBoost	0.5199	19.61	15.15
Holdout adsorbent	Screening-safe numeric-only LightGBM	0.5002	20.10	15.19
Holdout adsorbent	Full capacity XGBoost	0.4385	20.76	13.94
Holdout reference/source	Screening-safe numeric-only ExtraTrees	0.3054	23.54	18.67
Holdout reference/source	Full capacity LightGBM	0.2820	23.78	18.67
Holdout reference/source	Full capacity XGBoost	0.1678	25.66	19.53

These results show that removal-efficiency prediction is strongly source-dependent. The same data and model family can produce high R^2 under source-specific interpolation, moderate-to-good R^2 under random global valida-
245 tion, and lower but still informative R^2 under source-held-out validation when the applicability domain is restricted to sufficiently represented pH-concentration regimes.




Figure 3: Track B validation comparison. Row-wise validation, adsorbent-held-out validation, and reference-held-out validation represent increasingly strict generalization tasks.

3.4. Capacity-inclusive and capacity-free models answered different questions

250 The full capacity process model performed best for random row-wise validation and several source-specific models. This is expected because adsorption capacity carries strong information about adsorption performance. However, capacity can be outcome-adjacent to RE and may not be available before an experiment. Therefore, capacity-inclusive models should be framed as process-monitoring or
 255 descriptive models. Capacity-free models are more appropriate for pre-experiment screening and adsorbent-design guidance. In the restricted-domain adsorbent-held-out and reference-held-out tests, capacity-free numeric models were competitive or superior, suggesting that simpler descriptor sets can be more robust when transferring across material identities or source studies.

260 4. Discussion

4.1. *Why near-0.95 R^2 is achievable but validation-dependent*

The Track A results explain why many adsorption ML studies report R^2 values near or above 0.95. Within a single source, adsorbent family, pollutant system, or experimental protocol, the response surface can be learned accurately by tree-based
265 models. In this setting, high R^2 is scientifically plausible and not automatically invalid. The present results reproduce that behavior, with source-specific R^2 values of 0.960–0.992 for three major source groups.

However, Track B shows that this performance should not be interpreted as universal adsorption prediction. When sources are held out, the model must trans-
270 fer across experimental protocols, adsorbent synthesis methods, pollutant distributions, concentration ranges, and reporting conventions. Under that validation setting, performance fell below row-wise performance even after excluding the sparse neutral/low- C_0 regime. The result is consistent with recent adsorption-related work showing that random-split performance can be substantially higher
275 than leave-compound-out performance [9].

4.2. *Relation to existing adsorption ML literature*

The proposed framework complements rather than contradicts previous high-accuracy adsorption ML studies. Hafsa et al. demonstrated that tree-based algorithms can achieve high removal-efficiency prediction accuracy for
280 heavy-metal adsorption datasets under repeated cross-validation [1]. Wang et al. showed that XGBoost, SHAP, and partial-dependence analyses can identify interpretable drivers of heavy-metal removal by biochar [3]. Lu et al. showed that random forests can predict heavy-metal removal by chitosan-based flocculants

with R^2 above 0.93 [2]. Barkhordari et al. further demonstrated that interpretable
285 XGBoost models are relevant to JECE’s environmental-remediation scope by
predicting heavy-metal removal efficiency in electrokinetic remediation [5].
These studies establish that ML is useful for environmental treatment modeling.

The present study adds a validation-layer contribution: adsorption ML work-
flows should explicitly state whether they are predicting within a known source
290 or generalizing to unseen sources. This distinction is critical because high row-
wise R^2 can be overoptimistic when the deployment target is a new laboratory,
adsorbent synthesis route, pollutant matrix, or literature source. The broader ML
literature has identified leakage and validation mismatch as major reproducibil-
ity risks [7, 8], and adsorption datasets have several mechanisms by which such
295 mismatch can occur.

4.3. *Practical implications for adsorption model deployment*

For practical adsorption screening, a single universal RE model should be used
cautiously. A more defensible workflow is to define the pH-concentration appli-
cability domain before modeling, exclude or separately model poorly represented
300 regimes, use source-specific or laboratory-calibrated models for high-accuracy op-
timization within a known experimental system, use capacity-free screening mod-
els for preliminary adsorbent and operating-condition selection, and report source-
held-out validation when predictions are intended for new adsorbents, pollutants,
laboratories, or literature domains. This framing allows high- R^2 models to remain
305 useful without overstating their generality.

4.4. Limitations and future work

This study has four main limitations. First, the dataset does not contain explicit laboratory-origin labels for all records; bibliographic reference was therefore used as a source proxy. Future work should annotate records as in-house, external-
310 laboratory, or literature-derived and repeat validation with leave-laboratory-out splits. Second, the ion descriptor table should be expanded with fully cited descriptors such as electronegativity, hydration energy, softness/hardness, and likely aqueous speciation. Third, important adsorbent descriptors such as pH_{pzc} , pore volume, pore size, zeta potential, FTIR-derived functional groups, and cation-
315 exchange capacity are incomplete in the present dataset. Fourth, the neutral/low-concentration regime contained only 26 observations and was excluded from primary modeling; predictions in that regime should therefore be treated as outside the present applicability domain until more data are available. These limitations define clear routes for improving external predictive performance while preserving
320 the source-aware validation principle.

5. Conclusions

This study developed a source-aware and leakage-aware framework for adsorption removal-efficiency prediction. High R^2 values were achievable within individual adsorption sources, with source-specific XGBoost models reaching R^2 up
325 to 0.992. Because the neutral/low-concentration regime contained only 26 observations, primary global modeling excluded that sparse regime and used a 537-record restricted domain. Within this domain, random row-wise validation reached $R^2 = 0.806$ for the full capacity LightGBM process model, adsorbent-held-out validation reached $R^2 = 0.520$ for the capacity-free numeric XGBoost screening

330 model, and reference-held-out validation reached $R^2 = 0.305$ for the capacity-free
numeric ExtraTrees screening model. These findings show that high-accuracy in-
terpolation, applicability-domain definition, and cross-source generalization are
different scientific tasks. Future adsorption ML studies should report all three, sep-
arate capacity-inclusive process-monitoring models from capacity-free screening
335 models, and include source- or laboratory-held-out validation when deployment
beyond the training source is claimed.

CRedit authorship contribution statement

Yash Bisht: Conceptualization, Methodology, Software, Formal analysis, Vi-
sualization, Writing – original draft. Susmit Chitransh: Data curation, Investiga-
340 tion, Validation, Writing – review and editing.

Declaration of competing interest

The authors declare that they have no known competing financial interests or
personal relationships that could have appeared to influence the work reported in
this paper.

345 **Funding**

This research did not receive any specific grant from funding agencies in the
public, commercial, or not-for-profit sectors.

Data availability

The curated dataset, analysis scripts, and generated benchmark outputs are
350 available from the corresponding author upon reasonable request and will be de-

posited in a public repository before final publication.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During the preparation of this work, the authors used OpenAI's Codex to as-
355 sist with code organization, LaTeX formatting, and language polishing. After us-
ing this tool, the authors reviewed and edited the content as needed and take full
responsibility for the content of the manuscript.

Acknowledgements

The authors acknowledge the institutional support of the Indian Institute of
360 Technology Roorkee and the University of Cape Town.

References

- [1] N. Hafsa, S. Rushd, M. Al-Yaari, M. Rahman, A generalized method for modeling the adsorption of heavy metals with machine learning algorithms, Water 12 (2020) 3490. <https://doi.org/10.3390/w12123490>.
- 365 [2] C. Lu, Z. Xu, B. Dong, Y. Zhang, M. Wang, Y. Zeng, C. Zhang, Machine learning for the prediction of heavy metal removal by chitosan-based flocculants, Carbohydrate Polymers 285 (2022) 119240. <https://doi.org/10.1016/j.carbpol.2022.119240>.
- [3] C. Wang, Y. Zhao, Y. Gao, H. Chen, X. Li, B. Zhou, D. Fan, Z. Fang, J.
370 Liu, Interpretable machine learning for predicting heavy metal removal and

optimizing biochar characteristics, *Journal of Water Process Engineering* 68 (2024) 106484. <https://doi.org/10.1016/j.jwpe.2024.106484>.

- [4] Z.M. Yaseen, F.L. Alhalimi, Heavy metal adsorption efficiency prediction using biochar properties: a comparative analysis for ensemble machine learning models, *Scientific Reports* 15 (2025). <https://doi.org/10.1038/s41598-025-96271-5>.
375
- [5] M.S. Barkhordari, N. Zhou, K. Li, C. Qi, Interpretable machine learning for predicting heavy metal removal efficiency in electrokinetic soil remediation, *Journal of Environmental Chemical Engineering* 12 (2024) 114330. <https://doi.org/10.1016/j.jece.2024.114330>.
380
- [6] S.M. Lundberg, S.-I. Lee, A unified approach to interpreting model predictions, in: *Advances in Neural Information Processing Systems* 30, 2017, pp. 4765–4774.
- [7] S. Kapoor, A. Narayanan, Leakage and the reproducibility crisis in machine-learning-based science, *Patterns* 4 (2023) 100804. <https://doi.org/10.1016/j.patter.2023.100804>.
385
- [8] J.-J. Zhu, M. Yang, Z.J. Ren, Machine learning in environmental research: Common pitfalls and best practices, *Environmental Science & Technology* 57 (2023) 17671–17689. <https://doi.org/10.1021/acs.est.3c00026>.
- [9] H.V. Patel, J. Green, H. Park, S. Luster-Teasley Pass, R. Zhao, A hybrid response surface methodology and machine learning framework for quantifying effects of physicochemical parameters on PFAS distribution, *ACS*
390

ES&T Water (2026) Advance online publication. <https://doi.org/10.1021/acsestwater.5c01162>.

- 395 [10] T. Chen, C. Guestrin, XGBoost: A scalable tree boosting system, in: Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, ACM, San Francisco, CA, USA, 2016, pp. 785–794. <https://doi.org/10.1145/2939672.2939785>.
- [11] J.H. Friedman, Greedy function approximation: A gradient boosting machine, *Annals of Statistics* 29 (2001) 1189–1232. <https://doi.org/10.1214/aos/1013203451>.
400
- [12] L. Breiman, Random forests, *Machine Learning* 45 (2001) 5–32. <https://doi.org/10.1023/A:1010933404324>.
- [13] P. Geurts, D. Ernst, L. Wehenkel, Extremely randomized trees, *Machine Learning* 63 (2006) 3–42. <https://doi.org/10.1007/s10994-006-6226-1>.
405